



Geometric Deep Learning - Project Presentation

29/05/2019

Towards Sparse Hierarchical Graph Classifiers

Hierarchical Pooling for Graph Classification

Alessia Ruggeri

Introduction

The focus of this project is on the paper *Towards Sparse Hierarchical Graph Classifiers*, which present an interesting architecture for **hierarchical graph classification**.

The aim of the project is to analyze, implement and try to reproduce the results presented by the paper.

1. Problem formulation & Background

Problem formulation

The paper studies the problem of **graph classification**, related to the generalization of the pooling layer to graphs.

In order to preserve the topological information, a differentiable graph pooling layer is needed, which can generate **hierarchical representations** of graphs.

Background

Previous approaches to the **generalization of the pooling layer to graphs** through the aggregation of node representations:

- Global pooling layer

- Clusters which coarsen the graph in a hierarchical manner through a fixed and pre-defined cluster assignment

- Soft clustering assignments, learned in a differentiable way (DiffPool)

Background

The soft clustering assignment performed by **DiffPool** leads to state-of-the-art results on several graph classification benchmarks, but it requires a quadratic $O(kV^2)$ storage complexity to be executed.

This paper presents a **Hierarchical Pooling** included in an architecture able to replicate almost the same performance, requiring only $O(V + E)$ storage complexity.

2. Model

Model - Input graphs

The input graph is represented by:

Matrix of node features, $\mathbf{X} \in \mathbb{R}^{N \times F}$

Adjacency matrix, $\mathbf{A} \in \mathbb{R}^{N \times N}$

where N is the number of nodes in the graph and F the number of features for each node. A is binary and symmetric.

Model - Layers

In order to build the presented model architecture for graph classification, we need to define:

- Convolutional layer

- Pooling layer

- Readout layer

Model - Convolutional Layer

The architecture makes use of a **mean-pooling propagation rule**:

$$\text{MP}(\mathbf{X}, \mathbf{A}) = \sigma \left(\hat{\mathbf{D}}^{-1} \hat{\mathbf{A}} \mathbf{X} \boldsymbol{\Theta} + \underbrace{\mathbf{X} \boldsymbol{\Theta}'}_{\text{skip-connection}} \right)$$

where:

$\hat{\mathbf{A}} = \mathbf{A} + \mathbf{I}_N$ is the adjacency matrix with inserted self-loops

$\hat{\mathbf{D}}$ is the corresponding degree matrix, i.e. $\hat{D}_{ii} = \sum_j \hat{A}_{ij}$

σ is the rectified linear (ReLU) activation

$\boldsymbol{\Theta}, \boldsymbol{\Theta}' \in \mathbb{R}^{F \times F'}$ are learnable linear transformations applied to every node

Model - Pooling Layer

The hierarchical pooling layer reduces the graph with a **pooling ratio**, $k \in (0, 1]$, decreasing the number of nodes in a graph from N to $\lceil kN \rceil$. The choice of the nodes to drop is done based on a **projection score** against a learnable \vec{p} vector, which is also used as **gating value**.

Model - Pooling Layer


$$\vec{y} = \frac{\mathbf{X}\vec{p}}{\|\vec{p}\|} \quad \vec{i} = \text{top-}k(\vec{y}, k)$$

$$\tilde{\mathbf{X}} = \mathbf{X}(\vec{i}, :) \quad \mathbf{X}' = (\tilde{\mathbf{X}} \odot \tilde{y})$$

$$\tilde{y} = \tanh(\vec{y}(\vec{i}))$$

$$\mathbf{A}' = \mathbf{A}(\vec{i}, \vec{i})$$

where $\|\cdot\|$ is the L_2 norm and \odot is (broadcasted) element-wise multiplication.





Model - Readout Layer


The readout layer flatten the information about the input graph in a **fixed-size representation**. This operation is performed by **concatenating** the results of a global average pooling and a global max pooling at the end of each conv-pool block, and then applying a global sum pooling before submitting the information to an **MLP** to obtain final predictions.

Model - Readout Layer

$$\vec{s}^{(l)} = \frac{1}{N^{(l)}} \sum_{i=1}^{N^{(l)}} \vec{x}_i^{(l)} \parallel \max_{i=1}^N \vec{x}_i^{(l)} \quad \vec{s} = \sum_{l=1}^L \vec{s}^{(l)}$$


predictions = **MLP**(\vec{s})

where \parallel denotes the concatenation, $N^{(l)}$ the number of nodes of the graph and $\vec{x}_i^{(l)}$ the i -th node's feature vector.



Model - Pipeline

The presented graph neural network architecture comprises **three conv-pool blocks** (in the image they are only two).



3. Experiments

Datasets and Parameters

The graph neural network architecture of the paper has been evaluated on four benchmark tasks, with the following parameters:

Enzymes	128 feature layers, 0.0005 as learning rate, 100 epochs
Proteins	64 feature layers, 0.005 as learning rate, 40 epochs
D&D	64 feature layers, 0.0005 as learning rate, 20 epochs
Collab	128 feature layers, 0.0005 as learning rate, 30 epochs

For all the datasets, they set the pooling ratio $k = 0.8$ and used Adam optimizer.

For the evaluation of my implementation, I used only *Enzymes* and *Proteins* datasets.

My implementation

I replicated the same architecture described in the paper using **Tensorflow 2.0** on Google Colab and implementing all the layers from scratch.

In order to have a **baseline** for my hierarchical model, I also implemented a much simpler GCN model, which consists only of three convolutional layers, followed by a global max pooling layer just before the MLP for the classification.

I used the **same parameters** of the paper both for the hierarchical architecture and for the GCN model.


Results

The following table shows the classification accuracy percentages of the models:

Model	Datasets			
	<i>Enzymes</i>	<i>Proteins</i>	<i>DEED</i>	<i>Collab</i>
DiffPool-Det	58.33	75.62	75.47	82.13
DiffPool	64.23	78.10	81.15	75.50
Paper's model	64.17	75.46	78.59	74.54
My model	49.16	69.66	\	\
My baseline	55.83	63.09	\	\

Results

The following graph shows the GPU memory usage of the paper's method and DiffPool during training:



4. Conclusions

Replicated model

My implementation was **not able to reproduce the results** presented in the paper.

The obtained accuracy is very variable and hugely influenced by the parameters.

Probably the authors omitted crucial details in the paper.

Model	Datasets	
	<i>Enzymes</i>	<i>Proteins</i>
Paper's model	64.17	75.46
My model	49.16	69.66
My baseline	55.83	63.09

Thank you for your attention

Alessia Ruggeri